

Database Systems: Design, Implementation, and Management

Lesson 10

Database Performance-Tuning Concepts

- ▶ Goal of database performance is to execute queries as fast as possible
- ▶ **Database performance tuning**
 - Set of activities and procedures designed to reduce response time of database system
- ▶ All factors must operate at optimum level with minimal bottlenecks
- ▶ **Good database performance starts with good database design**

**TABLE
11.1**

General Guidelines for Better System Performance

	SYSTEM RESOURCES	CLIENT	SERVER
Hardware	CPU	The fastest possible dual-core CPU or higher	The fastest possible Multiple processors (quad-core technology)
	RAM	The maximum possible	The maximum possible
	Hard disk	Fast SATA/EIDE hard disk with sufficient free hard disk space	Multiple high-speed, high-capacity hard disks (SCSI/SATA/Firewire/Fibre Channel) in RAID configuration
	Network	High-speed connection	High-speed connection
Software	Operating system	Fine-tuned for best client application performance	Fine-tuned for best server application performance
	Network	Fine-tuned for best throughput	Fine-tuned for best throughput
	Application	Optimize SQL in client application	Optimize DBMS server for best performance

Performance Tuning: Client and Server

▶ Client side

- Generate SQL query that returns correct answer in least amount of time
 - Using minimum amount of resources at server
- **SQL performance tuning**

▶ Server side

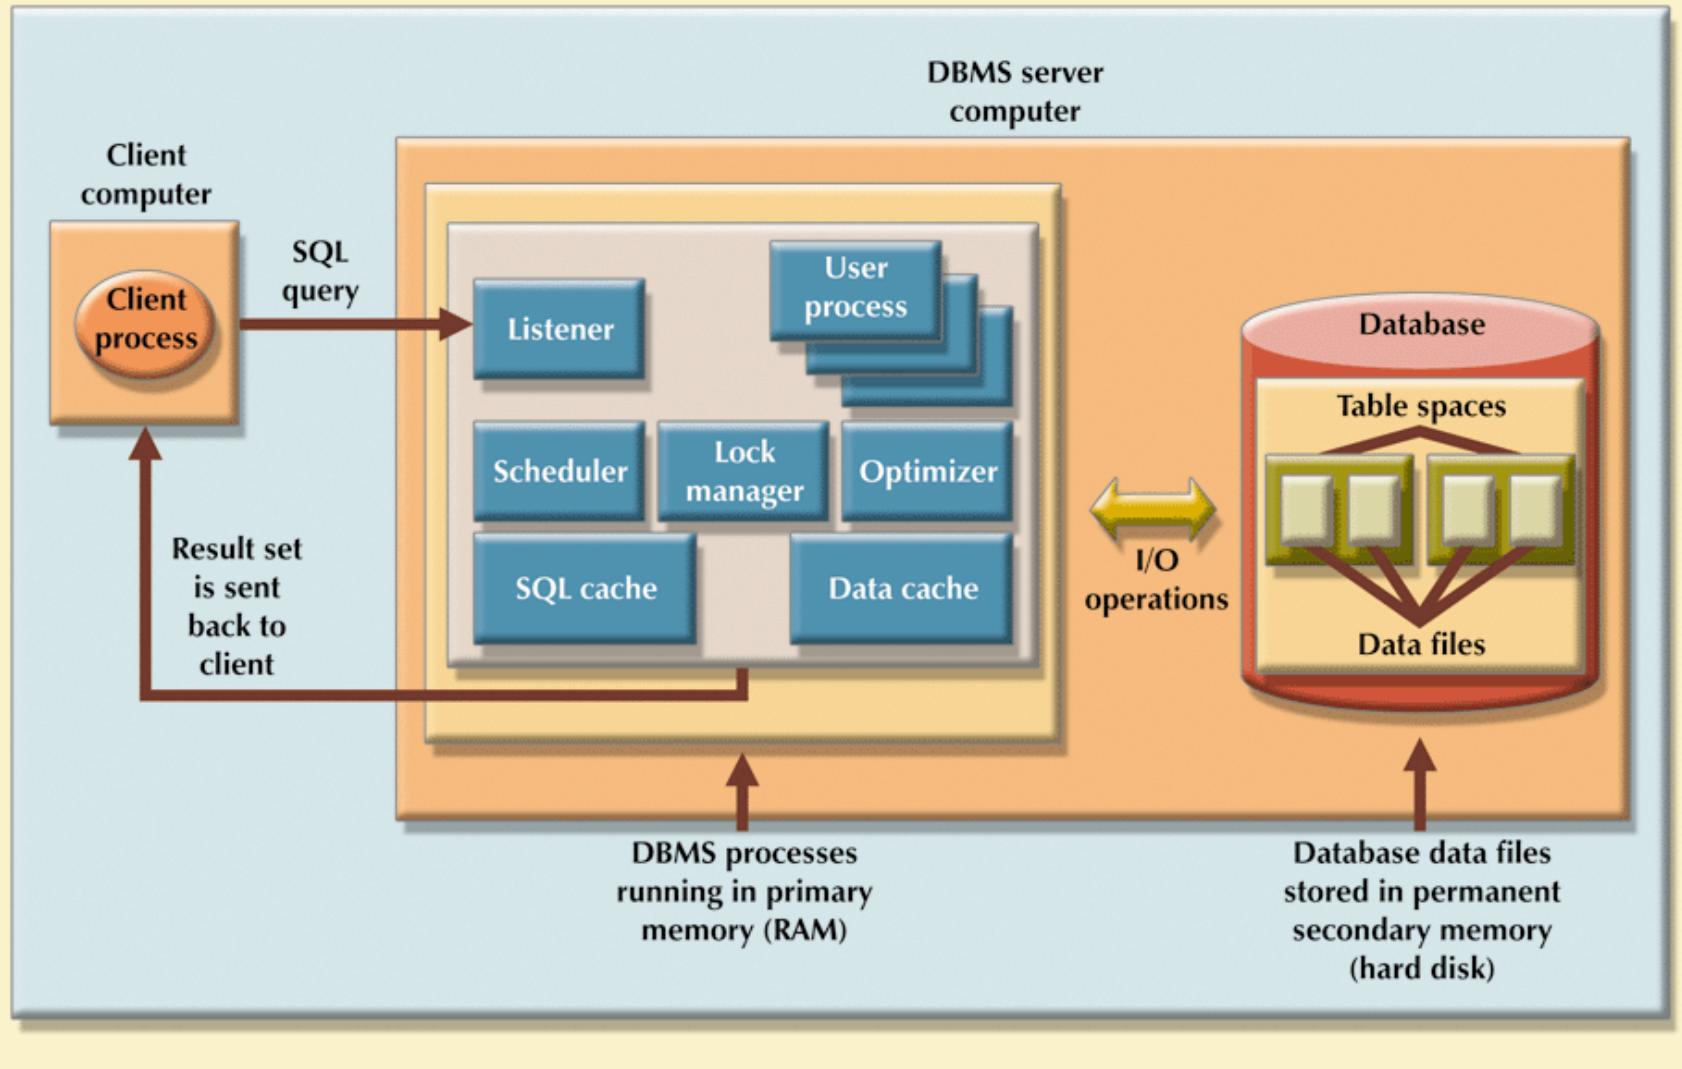
- DBMS environment configured to respond to clients' requests as fast as possible
 - Optimum use of existing resources
- **DBMS performance tuning**

DBMS Architecture

- ▶ All data in database are stored in data files
- ▶ **Data files**
 - Automatically expand in predefined increments known as **extends**
 - Grouped in **file groups** or **table spaces**
- ▶ **Table space or file group**
 - Logical grouping of several data files that store data with similar characteristics

FIGURE
11.1

Basic DBMS architecture



DBMS Architecture (cont'd.)

- ▶ **Data cache or buffer cache:** shared, reserved memory area
 - Stores most recently accessed data blocks in RAM
- ▶ **SQL cache or procedure cache:** stores most recently executed SQL statements
 - Also PL/SQL procedures, including triggers and functions
- ▶ DBMS retrieves data from permanent storage and places it in RAM

DBMS Architecture (cont'd.)

- ▶ **Input/output request:** low-level data access operation to/from computer devices
- ▶ Data cache is faster than data in data files
 - DBMS does not wait for hard disk to retrieve data
- ▶ Majority of performance-tuning activities focus on minimizing I/O operations
- ▶ Typical DBMS processes:
 - Listener, user, scheduler, lock manager, optimizer

Database Statistics

- ▶ Measurements about database objects and available resources:
 - Tables
 - Indexes
 - Number of processors used
 - Processor speed
 - Temporary space available

Database Statistics (cont'd.)

- ▶ Make critical decisions about improving query processing efficiency
- ▶ Can be gathered manually by DBA or automatically by DBMS

**TABLE
11.2**

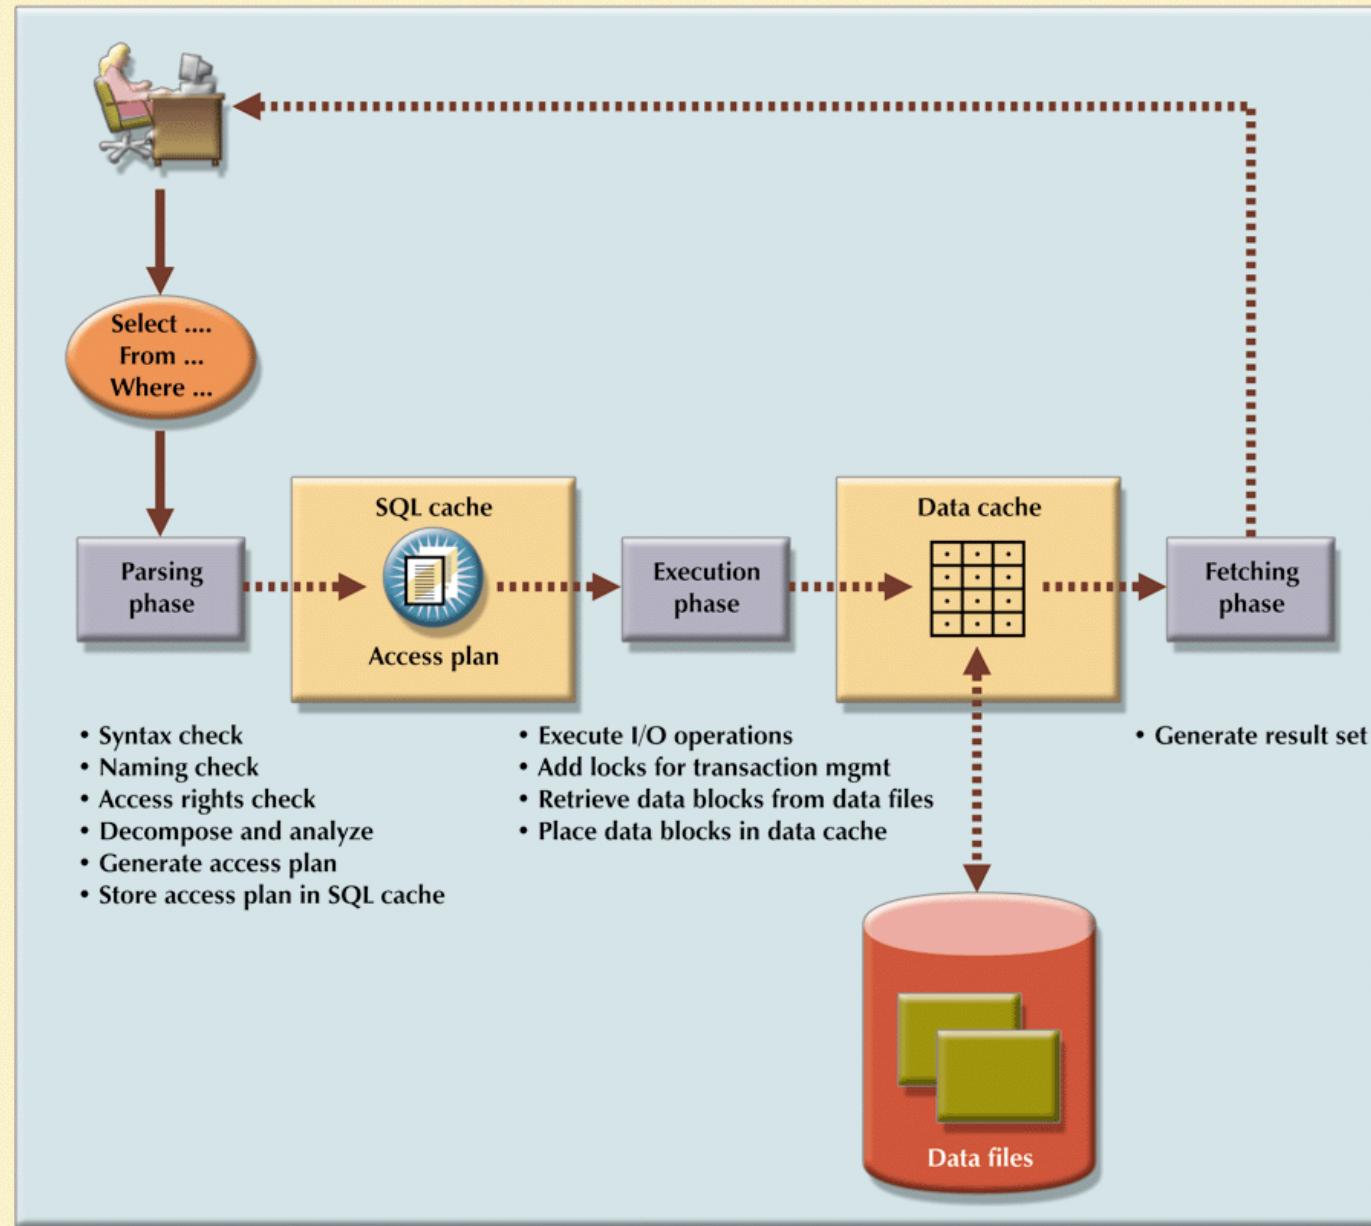
Sample Database Statistics Measurements

DATABASE OBJECT	SAMPLE MEASUREMENTS
Tables	Number of rows, number of disk blocks used, row length, number of columns in each row, number of distinct values in each column, maximum value in each column, minimum value in each column, and columns that have indexes
Indexes	Number and name of columns in the index key, number of key values in the index, number of distinct key values in the index key, histogram of key values in an index, and number of disk pages used by the index
Environment Resources	Logical and physical disk block size, location and size of data files, and number of extends per data file

Query Processing

- ▶ DBMS processes queries in three phases
 - Parsing
 - DBMS parses the query and chooses the most efficient access/execution plan
 - Execution
 - DBMS executes the query using chosen execution plan
 - Fetching
 - DBMS fetches the data and sends the result back to the client

FIGURE 11.2 Query processing



Indexes and Query Optimization

- ▶ Indexes
 - Crucial in speeding up data access
 - Facilitate searching, sorting, and using aggregate functions as well as join operations
 - Ordered set of values that contains index key and pointers
- ▶ More efficient to use index to access table than to scan all rows in table sequentially

Indexes and Query Optimization (cont'd.)

- ▶ **Data sparsity:** number of different values a column could possibly have
- ▶ Indexes implemented using:
 - Hash indexes
 - B-tree indexes
 - Bitmap indexes
- ▶ DBMSs determine best type of index to use

**FIGURE
11.3**

Index representation for the CUSTOMER table

STATE_NDX INDEX

Key	Row
AZ	2
....
....
FL	1
FL	7
FL	8
FL	13245
FL	14786
....
....

CUSTOMER TABLE (14,786 rows)									
	Row ID	CUS_CODE	CUS_LNAME	CUS_FNAME	CUS_INITIAL	CUS_AREACODE	CUS_PHONE	CUS_STATE	CUS_BALANCE
	1	10010	Ramas	Alfred	A	615	844-2573	FL	\$0.00
	2	10011	Dunne	Leona	K	713	894-1238	AZ	\$0.00
	3	10012	Smith	Kathy	W	615	894-2285	TX	\$345.86
	4	10013	Olowski	Paul	F	615	894-2180	AZ	\$536.75
	5	10014	Orlando	Myron		615	222-1672	NY	\$0.00
	6	10015	O'Brian	Amy	B	713	442-3381	NY	\$0.00
	7	10016	Brown	James	G	615	297-1228	FL	\$221.19
	8	10017	Williams	George		615	290-2556	FL	\$768.93
	9	10018	Farriss	Anne	G	713	382-7185	TX	\$216.55
	10	10019	Smith	Olette	K	615	297-3809	AZ	\$0.00

	13245	23120	Veron	George	D	415	231-9872	FL	\$675.00

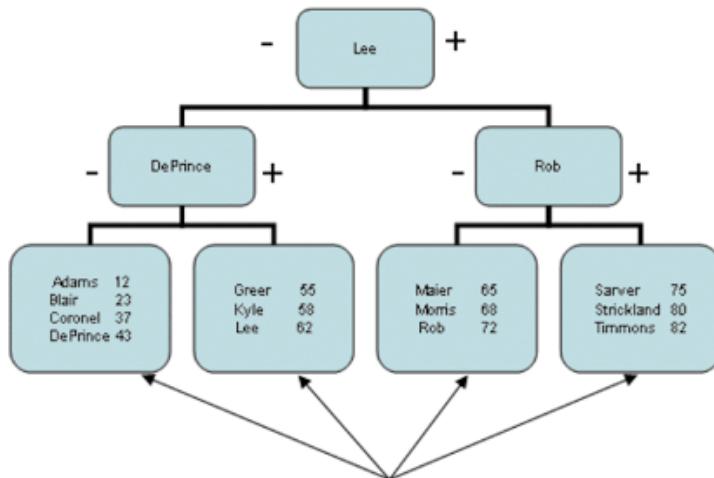
	14786	24560	Suarez	Victor		435	342-9876	FL	\$342.00

FIGURE
11.4

B-tree and bitmap index representation

B-Tree index is used in columns with high data sparsity – that is, columns with many different values relative to the total number of rows.

B-tree Index
On CUS_LNAME



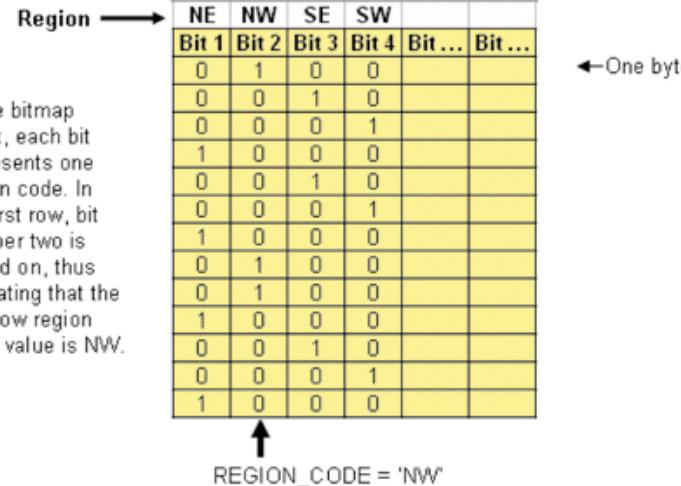
Leaf objects contain index key and pointers to rows in table.
 Access to any row using the index will take the same number of I/O accesses. In this example, it would take four I/O accesses to access any given table row using the index: One for each index tree level (root, branch, leaf object) plus access to data row using the pointer.

CUSTOMER TABLE

CUS_ID	CUS_LNAME	CUS_FNAME	CUS_PHONE	REGION_CODE
12	Adams	Charlie	4533	NW
23	Blair	Robert	5426	SE
37	Coronel	Carlos	2358	SW
43	DePrince	Albert	6543	NE
55	Greer	Tim	2764	SE
58	Kyle	Ruben	2453	SW
62	Lee	John	7895	NE
65	Maier	Jerry	7689	NW
68	Morris	Steve	4568	NW
72	Rob	Pete	8123	NE
75	Sarver	Lee	8193	SE
80	Strickland	Tomas	3129	SW
82	Timmons	Douglas	3499	NE

Bitmap index is used in columns with low data sparsity – that is, columns with few different values relative to the total number of rows.

Bitmap Index
On REGION_CODE



Each byte in the bitmap index represents one row of the table data. Bitmap indexes are very efficient with searches. For example, to find all customers in the NW region, the DBMS will return all rows with bit number two turned on.

SQL Performance Tuning

- ▶ Evaluated from client perspective
 - Most current relational DBMSs perform automatic query optimization at the server end
 - Most SQL performance optimization techniques are DBMS specific
 - Rarely portable
- ▶ Majority of performance problems are related to poorly written SQL code
- ▶ Carefully written query usually outperforms a poorly written query

Index Selectivity

- ▶ Indexes are used when:
 - Indexed column appears by itself in search criteria of WHERE or HAVING clause
 - Indexed column appears by itself in GROUP BY or ORDER BY clause
 - MAX or MIN function is applied to indexed column
 - Data sparsity on indexed column is high
- ▶ Measure of how likely an index will be used

Index Selectivity (cont'd.)

- ▶ General guidelines for indexes:
 - Create indexes for each attribute in WHERE, HAVING, ORDER BY, or GROUP BY clause
 - Do not use in small tables or tables with low sparsity
 - Declare primary and foreign keys so optimizer can use indexes in join operations
 - Declare indexes in join columns other than PK/FK

Conditional Expressions

- ▶ Normally expressed within WHERE or HAVING clauses of SQL statement
- ▶ Restricts output of query to only rows matching conditional criteria

Conditional Expressions (cont'd.)

- ▶ Common practices for efficient SQL:
 - Use simple columns or literals in conditionals
 - Numeric field comparisons are faster
 - Equality comparisons are faster than inequality
 - Transform conditional expressions to use literals
 - Write equality conditions first
 - AND: use condition most likely to be false first
 - OR: use condition most likely to be true first
 - Avoid NOT

Query Formulation

- ▶ Identify what columns and computations are required
- ▶ Identify source tables
- ▶ Determine how to join tables
- ▶ Determine what selection criteria is needed
- ▶ Determine in what order to display output

DBMS Performance Tuning

- ▶ Includes managing DBMS processes in primary memory and structures in physical storage
- ▶ DBMS performance tuning at server end focuses on setting parameters used for:
 - Data cache
 - SQL cache
 - Sort cache
 - Optimizer mode

DBMS Performance Tuning (cont'd.)

- ▶ Some general recommendations for creation of databases:
 - Use RAID (Redundant Array of Independent Disks) to provide balance between performance and fault tolerance
 - Minimize disk contention
 - Put high-usage tables in their own table spaces
 - Assign separate data files in separate storage volumes for indexes, system, high-usage tables

DBMS Performance Tuning (cont'd.)

- ▶ Some general recommendations for creation of databases (cont'd):
 - Take advantage of table storage organizations in database
 - Partition tables based on usage
 - Use denormalized tables where appropriate
 - Store computed and aggregate attributes in tables